

APP RATING PREDICTION

Objective: Make a model to predict the app rating, with other information about the app provided.

Problem Statement:

Google Play Store team is about to launch a new feature wherein, certain apps that are promising, are boosted in visibility. The boost will manifest in multiple ways including higher priority in recommendations sections ("Similar apps", "You might also like", "New and updated games"). These will also get a boost in search results visibility. This feature will help bring more attention to newer apps that have the potential.

Steps to perform:

1. Load the data file using pandas

```
import pandas as pd
import numpy as np
import seaborn as sns
df=pd.read_csv("googleplaystore.csv")
df.head()
```

Rating \	App	Category
0	Photo Editor & Candy Camera & Grid & ScrapBook	ART_AND_DESIGN
4.1		
1	Coloring book moana	ART_AND_DESIGN
3.9		
2	U Launcher Lite – FREE Live Cool Themes, Hide ...	ART_AND_DESIGN
4.7		
3	Sketch - Draw & Paint	ART_AND_DESIGN
4.5		
4	Pixel Draw - Number Art Coloring Book	ART_AND_DESIGN
4.3		

	Reviews	Size	Installs	Type	Price	Content	Rating \
0	159	19M	10,000+	Free	0		Everyone
1	967	14M	500,000+	Free	0		Everyone
2	87510	8.7M	5,000,000+	Free	0		Everyone
3	215644	25M	50,000,000+	Free	0		Teen
4	967	2.8M	100,000+	Free	0		Everyone

	Genres	Last Updated	Current Ver \
0	Art & Design	January 7, 2018	1.0.0
1	Art & Design;Pretend Play	January 15, 2018	2.0.0
2	Art & Design	August 1, 2018	1.2.4
3	Art & Design	June 8, 2018	Varies with device

4 Art & Design;Creativity June 20, 2018 1.1

Android Ver
0 4.0.3 and up
1 4.0.3 and up
2 4.0.3 and up
3 4.2 and up
4 4.4 and up

df.tail()

App
Category \
10836 Sya9a Maroc - FR
FAMILY
10837 Fr. Mike Schmitz Audio Teachings
FAMILY
10838 Parkinson Exercices FR
MEDICAL
10839 The SCP Foundation DB fr nn5n
BOOKS_AND_REFERENCE
10840 iHoroscope - 2018 Daily Horoscope & Astrology
LIFESTYLE

	Rating	Reviews	Size	Installs	Type	Price	\
10836	4.5	38	53M	5,000+	Free	0	
10837	5.0	4	3.6M	100+	Free	0	
10838	NaN	3	9.5M	1,000+	Free	0	
10839	4.5	114	Varies with device	1,000+	Free	0	
10840	4.5	398307	19M	10,000,000+	Free	0	

	Content Rating	Genres	Last Updated
Current Ver \			
10836	Everyone	Education	July 25, 2017
1.48			
10837	Everyone	Education	July 6, 2018
1.0			
10838	Everyone	Medical	January 20, 2017
1.0			
10839	Mature 17+	Books & Reference	January 19, 2015
device			Varies with
10840	Everyone	Lifestyle	July 25, 2018
device			Varies with

Android Ver
10836 4.1 and up
10837 4.1 and up
10838 2.2 and up
10839 Varies with device
10840 Varies with device

2. Check for null values in the data. Get the number of null values for each column.

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10841 entries, 0 to 10840
Data columns (total 13 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   App                   10841 non-null  object
 1   Category              10841 non-null  object
 2   Rating                9367 non-null   float64
 3   Reviews               10841 non-null  object
 4   Size                  10841 non-null  object
 5   Installs              10841 non-null  object
 6   Type                  10840 non-null  object
 7   Price                 10841 non-null  object
 8   Content Rating        10840 non-null  object
 9   Genres                10841 non-null  object
10   Last Updated          10841 non-null  object
11   Current Ver           10833 non-null  object
12   Android Ver           10838 non-null  object
dtypes: float64(1), object(12)
memory usage: 1.1+ MB
```

```
df.isnull().sum()
```

```
App                0
Category           0
Rating            1474
Reviews            0
Size               0
Installs           0
Type               1
Price              0
Content Rating     1
Genres             0
Last Updated       0
Current Ver        8
Android Ver        3
dtype: int64
```

3. Drop records with nulls in any of the columns.

```
df.dropna(inplace=True)
df.isnull().sum()
```

```
App                0
Category           0
Rating             0
Reviews            0
Size               0
```

```

Installs          0
Type              0
Price             0
Content Rating    0
Genres            0
Last Updated      0
Current Ver       0
Android Ver       0
dtype: int64

```

4. Variables seem to have incorrect type and inconsistent formatting. You need to fix them:

4.1. Size column has sizes in Kb as well as Mb. To analyze, you'll need to convert these to numeric.

```

df['Size'].unique()

array(['19M', '14M', '8.7M', '25M', '2.8M', '5.6M', '29M', '33M',
      '3.1M', '28M', '12M', '20M', '21M', '37M', '5.5M', '17M', '39M', '31M',
      '4.2M', '23M', '6.0M', '6.1M', '4.6M', '9.2M', '5.2M', '11M',
      '24M', 'Varies with device', '9.4M', '15M', '10M', '1.2M',
      '26M', '8.0M', '7.9M', '56M', '57M', '35M', '54M', '201k', '3.6M',
      '5.7M', '8.6M', '2.4M', '27M', '2.7M', '2.5M', '7.0M', '16M', '3.4M',
      '8.9M', '3.9M', '2.9M', '38M', '32M', '5.4M', '18M', '1.1M',
      '2.2M', '4.5M', '9.8M', '52M', '9.0M', '6.7M', '30M', '2.6M',
      '7.1M', '22M', '6.4M', '3.2M', '8.2M', '4.9M', '9.5M', '5.0M',
      '5.9M', '13M', '73M', '6.8M', '3.5M', '4.0M', '2.3M', '2.1M',
      '42M', '9.1M', '55M', '23k', '7.3M', '6.5M', '1.5M', '7.5M',
      '51M', '41M', '48M', '8.5M', '46M', '8.3M', '4.3M', '4.7M', '3.3M',
      '40M', '7.8M', '8.8M', '6.6M', '5.1M', '61M', '66M', '79k', '8.4M',
      '3.7M', '118k', '44M', '695k', '1.6M', '6.2M', '53M', '1.4M',
      '3.0M', '7.2M', '5.8M', '3.8M', '9.6M', '45M', '63M', '49M',
      '77M', '4.4M', '70M', '9.3M', '8.1M', '36M', '6.9M', '7.4M', '84M',
      '97M', '2.0M', '1.9M', '1.8M', '5.3M', '47M', '556k', '526k', '76M',
      '7.6M', '59M', '9.7M', '78M', '72M', '43M', '7.7M', '6.3M',
      '334k', '93M', '65M', '79M', '100M', '58M', '50M', '68M', '64M', '34M',
      '67M', '60M', '94M', '9.9M', '232k', '99M', '624k', '95M',
      '8.5k', '41k', '292k', '80M', '1.7M', '10.0M', '74M', '62M', '69M',
      '75M', '98M', '85M', '82M', '96M', '87M', '71M', '86M', '91M', '81M',
      '92M', '83M', '88M', '704k', '862k', '899k', '378k', '4.8M',
      '266k', '375k', '1.3M', '975k', '980k', '4.1M', '89M', '696k',

```

```

'544k', '525k', '920k', '779k', '853k', '720k', '713k', '772k',
'318k', '58k', '241k', '196k', '857k', '51k', '953k', '865k',
'251k', '930k', '540k', '313k', '746k', '203k', '26k', '314k',
'239k', '371k', '220k', '730k', '756k', '91k', '293k', '17k',
'74k', '14k', '317k', '78k', '924k', '818k', '81k', '939k',
'169k',
'45k', '965k', '90M', '545k', '61k', '283k', '655k', '714k',
'93k',
'872k', '121k', '322k', '976k', '206k', '954k', '444k', '717k',
'210k', '609k', '308k', '306k', '175k', '350k', '383k', '454k',
'1.0M', '70k', '812k', '442k', '842k', '417k', '412k', '459k',
'478k', '335k', '782k', '721k', '430k', '429k', '192k', '460k',
'728k', '496k', '816k', '414k', '506k', '887k', '613k', '778k',
'683k', '592k', '186k', '840k', '647k', '373k', '437k', '598k',
'716k', '585k', '982k', '219k', '55k', '323k', '691k', '511k',
'951k', '963k', '25k', '554k', '351k', '27k', '82k', '208k',
'551k', '29k', '103k', '116k', '153k', '209k', '499k', '173k',
'597k', '809k', '122k', '411k', '400k', '801k', '787k', '50k',
'643k', '986k', '516k', '837k', '780k', '20k', '498k', '600k',
'656k', '221k', '228k', '176k', '34k', '259k', '164k', '458k',
'629k', '28k', '288k', '775k', '785k', '636k', '916k', '994k',
'309k', '485k', '914k', '903k', '608k', '500k', '54k', '562k',
'847k', '948k', '811k', '270k', '48k', '523k', '784k', '280k',
'24k', '892k', '154k', '18k', '33k', '860k', '364k', '387k',
'626k', '161k', '879k', '39k', '170k', '141k', '160k', '144k',
'143k', '190k', '376k', '193k', '473k', '246k', '73k', '253k',
'957k', '420k', '72k', '404k', '470k', '226k', '240k', '89k',
'234k', '257k', '861k', '467k', '676k', '552k', '582k',
'619k'],
dtype=object)

```

```

def change(Size):
    if 'M' in Size:
        x=Size[:-1]
        x=float(x)*1000
        return x

    elif 'K' in Size:
        x=Size[:-1]
        x=float(x)
        return x

    else: return None

```

```

df.Size=df.Size.map(change)
df.Size.value_counts()

```

```

14000.0    165
12000.0    161
11000.0    159

```

```

15000.0    159
13000.0    157
...
89000.0     9
84000.0     9
86000.0     8
90000.0     5
1000.0      4
Name: Size, Length: 181, dtype: int64

```

```
df['Size'].unique()
```

```

array([ 19000., 14000.,  8700., 25000.,  2800.,  5600., 29000.,
        33000.,  3100., 28000., 12000., 20000., 21000., 37000.,
         5500., 17000., 39000., 31000.,  4200., 23000.,  6000.,
         6100.,  4600.,  9200.,  5200., 11000., 24000.,    nan,
         9400., 15000., 10000.,  1200., 26000.,  8000.,  7900.,
        56000., 57000., 35000., 54000.,  3600.,  5700.,  8600.,
         2400., 27000.,  2700.,  2500.,  7000., 16000.,  3400.,
         8900.,  3900.,  2900., 38000., 32000.,  5400., 18000.,
         1100.,  2200.,  4500.,  9800., 52000.,  9000.,  6700.,
        30000.,  2600.,  7100., 22000.,  6400.,  3200.,  8200.,
         4900.,  9500.,  5000.,  5900., 13000., 73000.,  6800.,
         3500.,  4000.,  2300.,  2100., 42000.,  9100., 55000.,
         7300.,  6500.,  1500.,  7500., 51000., 41000., 48000.,
         8500., 46000.,  8300.,  4300.,  4700.,  3300., 40000.,
         7800.,  8800.,  6600.,  5100., 61000., 66000.,  8400.,
         3700., 44000.,  1600.,  6200., 53000.,  1400.,  3000.,
         7200.,  5800.,  3800.,  9600., 45000., 63000., 49000.,
        77000.,  4400., 70000.,  9300.,  8100., 36000.,  6900.,
         7400., 84000., 97000.,  2000.,  1900.,  1800.,  5300.,
        47000., 76000.,  7600., 59000.,  9700., 78000., 72000.,
        43000.,  7700.,  6300., 93000., 65000., 79000., 100000.,
        58000., 50000., 68000., 64000., 34000., 67000., 60000.,
        94000.,  9900., 99000., 95000., 80000.,  1700., 74000.,
        62000., 69000., 75000., 98000., 85000., 82000., 96000.,
        87000., 71000., 86000., 91000., 81000., 92000., 83000.,
        88000.,  4800.,  1300.,  4100., 89000., 90000.,  1000.])

```

4.2.Reviews is a numeric field that is loaded as a string field. Convert it to numeric (int/float).

```
df.Reviews.dtype
```

```
dtype('O')
```

```
df['Reviews']=df.Reviews.replace('3.0M',3000000.0)
```

```
df['Reviews']=df['Reviews'].astype('int')
```

```
df.Reviews.dtype
```

```
dtype('int64')
```

4.3. Installs field is currently stored as string and has values like 1,000,000+.

```
df['Installs'].unique()

array(['10,000+', '500,000+', '5,000,000+', '50,000,000+', '100,000+',
      '50,000+', '1,000,000+', '10,000,000+', '5,000+',
      '100,000,000+',
      '1,000,000,000+', '1,000+', '500,000,000+', '100+', '500+',
      '10+',
      '5+', '50+', '1+'], dtype=object)
```

4.3.1. Treat 1,000,000+ as 1,000,000

```
df['Installs']=df.Installs.str.replace("+","")
df.Installs=df.Installs.str.replace('Free','0')

df['Installs'].unique()

array(['10,000', '500,000', '5,000,000', '50,000,000', '100,000',
      '50,000', '1,000,000', '10,000,000', '5,000', '100,000,000',
      '1,000,000,000', '1,000', '500,000,000', '100', '500', '10',
      '5',
      '50', '1'], dtype=object)
```

4.3.2. remove '+', ',' from the field, convert it to integer

```
df['Installs']=df.Installs.str.replace(",","",")

df.Installs=pd.to_numeric(df.Installs)

df.Installs.dtype

dtype('int64')
```

4.4. Price field is a string and has \$ symbol. Remove '\$' sign, and convert it to numeric.

```
df.Price=df.Price.str.replace('$','')
df.Price=df.Price.str.replace('Everyone','0')

df.Price=pd.to_numeric(df.Price)

df.Price.unique()

array([ 0. ,  4.99,  3.99,  6.99,  7.99,  5.99,  2.99,  3.49,
        1.99,  9.99,  7.49,  0.99,  9. ,  5.49, 10. , 24.99,
       11.99, 79.99, 16.99, 14.99, 29.99, 12.99,  2.49, 10.99,
        1.5 , 19.99, 15.99, 33.99, 39.99,  3.95,  4.49,  1.7 ,
        8.99,  1.49,  3.88, 399.99, 17.99, 400. ,  3.02,  1.76,
        4.84,  4.77,  1.61,  2.5 ,  1.59,  6.49,  1.29, 299.99,
       379.99, 37.99, 18.99, 389.99,  8.49,  1.75, 14. ,  2. ,
        3.08,  2.59, 19.4 ,  3.9 ,  4.59, 15.46,  3.04, 13.99,
```

```
4.29, 3.28, 4.6 , 1. , 2.95, 2.9 , 1.97, 2.56,  
1.2 ])
```

```
df.Price.dtype
```

```
dtype('float64')
```

5.Sanity checks:

5.1.Average rating should be between 1 and 5 as only these values are allowed on the play store. Drop the rows that have a value outside this range.

```
df=df[(df.Rating>=1)&(df.Rating<=5)]
```

```
df.Rating.value_counts()
```

```
4.4    1108  
4.3    1076  
4.5    1037  
4.2     951  
4.6     823  
4.1     707  
4.0     567  
4.7     499  
3.9     386  
3.8     303  
5.0     274  
3.7     239  
4.8     234  
3.6     174  
3.5     163  
3.4     128  
3.3     102  
4.9      87  
3.0      83  
3.1      69  
3.2      63  
2.9      45  
2.8      42  
2.6      25  
2.7      25  
2.5      21  
2.3      20  
2.4      19  
1.0      16  
2.2      14  
1.9      13  
2.0      12  
1.7       8  
2.1       8  
1.8       8  
1.6       4
```



```
1.4      3
1.5      3
1.2      1
Name: Rating, dtype: int64
```

```
df['Rating'].unique()
```

```
array([4.1, 3.9, 4.7, 4.5, 4.3, 4.4, 3.8, 4.2, 4.6, 4. , 4.8, 4.9,
3.6,
      3.7, 3.2, 3.3, 3.4, 3.5, 3.1, 5. , 2.6, 3. , 1.9, 2.5, 2.8,
2.7,
      1. , 2.9, 2.3, 2.2, 1.7, 2. , 1.8, 2.4, 1.6, 2.1, 1.4, 1.5,
1.2])
```

5.2.Reviews should not be more than installs as only those who installed can review the app. If there are any such records, drop them.

```
df.Reviews
```

```
0          159
1          967
2         87510
3        215644
4          967
...
10834         7
10836        38
10837         4
10839        114
10840       398307
Name: Reviews, Length: 9360, dtype: int64
```

```
df.Installs
```

```
0          10000
1         500000
2        5000000
3       50000000
4        100000
...
10834         500
10836        5000
10837         100
10839        1000
10840       10000000
Name: Installs, Length: 9360, dtype: int64
```

```
len(df.Installs)
```

```
9360
```

```
len(df.Rating)
```

9360

5.3. For free apps (type = "Free"), the price should not be >0. Drop any such rows.

```
df.Type.value_counts()
```

```
Free      8715
```

```
Paid       645
```

```
Name: Type, dtype: int64
```

```
index_free_price_0=df.index[((df.Type=='Free')&(df.Price>0))]
```

```
if len(index_free_price_0)>0:
```

```
    print("Dropping such values",index_free_price_0)
```

```
    df.drop(index_free_price_0,axis=0,inplace=True)
```

```
else:
```

```
    print("There has no apps price>0")
```

There has no apps price>0

1. Performing univariate analysis:

5.1.Boxplot for Price

5.1.Q.Are there any outliers? Think about the price of usual apps on Play Store.


```
sns.boxplot(df['Price'])
```

```
/usr/local/lib/python3.7/site-packages/seaborn/_decorators.py:43:
```

```
FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
```

```
FutureWarning
```

```
<AxesSubplot:xlabel='Price'>
```



5.1.ANS. most of the price values are located in the range of 50. greater than 100 may be considered as outliers

```
import matplotlib.pyplot as plt
```

```
import statistics as stc
```

standard deviation of price

```
p_std=stc.stdev(df.Price)
```

```
print(p_std)
```

```
15.82164024735431
```

mean of price

```
p_mean=stc.mean(df.Price)
```

```
print(p_mean)
```

```
0.9612788461538462
```

price upper limit

```
price_up_lim=p_mean+3*p_std
```

```
print(price_up_lim)
```

```
48.426199588216775
```

```
len(df[df.Price>price_up_lim])
```

```
17
```

price lower limit

```
price_low_lim=p_mean-3*p_std  
print(price_low_lim)
```

-46.503641895909084

```
len(df[df.Price<price_low_lim])
```

0

here i got 17 outliers

5.2.Boxplot for Reviews


5.2.Q.Are there any apps with very high number of reviews? Do the values seem right?

```
sns.boxplot(df['Reviews'])
```

```
/usr/local/lib/python3.7/site-packages/seaborn/_decorators.py:43:  
FutureWarning: Pass the following variable as a keyword arg: x. From  
version 0.12, the only valid positional argument will be `data`, and  
passing other arguments without an explicit keyword will result in an  
error or misinterpretation.
```

FutureWarning

<AxesSubplot:xlabel='Reviews'>



5.2.ANS.here we can find out some outliers.

standard deviation reviews

```
review_std=stc.stdev(df.Reviews)
print(review_std)
```

```
3145023.255620224
```

mean of review

```
review_mean=stc.mean(df.Reviews)
print(review_mean)
```

```
514376.7052350427
```

review uper limit

```
review_up_lim=review_mean+3*review_std
print(review_up_lim)
```

```
9949446.472095715
```

```
len(df[df.Reviews>review_up_lim])
```

```
92
```

review lower limt

```
review_low_lim=review_mean-3*review_std
print(review_low_lim)
```

```
-8920693.061625628
```

```
len(df[df.Reviews<review_low_lim])
```

```
0
```

review column has 60 upper outlier.

removeing outlier

```
df.drop(df.index[(df.Reviews>review_up_lim)],inplace=True)
```

```
len(df[df.Reviews>review_up_lim])
```


```
0
```

5.3.Histogram for Rating

5.3.Q.How are the ratings distributed? Is it more toward higher ratings?

```
sns.displot(df.Rating)
```

```
<seaborn.axisgrid.FacetGrid at 0x7fa239ce70d0>
```



5.3.ans.this graph showing that 4.0 to 4.7 high variation(peak).it is a left skew plot.

define mean and standard deviation

```
rating_mean=np.mean(df.Rating)
rating_mean
```

```
4.189846784635309
```

```
rating_std=np.std(df.Rating)
rating_std
```

```
0.5170834870588434
```

define uper and lower outliers

```
rating_up_lim=rating_mean+3*rating_std
print(rating_up_lim)
```

```
5.741097245811839
```

```
len(df[df.Rating>rating_up_lim])
```

```
0
```

```
rating_low_lim=rating_mean-3*rating_std  
rating_low_lim
```

```
2.6385963234587786
```

```
len(df[df.Rating<rating_low_lim])
```

```
175
```

removing outliers


```
df.drop(df.index[df.Rating<rating_low_lim],inplace=True)  
len(df[df.Rating<rating_low_lim])
```

```
0
```

5.4.Histogram for Size


```
sns.displot(df.Size)
```

```
<seaborn.axisgrid.FacetGrid at 0x7fa239912f90>
```



```
sns.histplot(x='Size',data=df,log_scale=True)
```

```
<AxesSubplot:xlabel='Size', ylabel='Count'>
```



```
size_mean=stc.mean(df.Size)
print(size_mean)
```

nan

```
size_std=stc.stdev(df.Size)
print(size_std)
```

nan

6.Outlier treatment:

6.1.Price: From the box plot, it seems like there are some apps with very high price. A price of \$200 for an application on the Play Store is very high and suspicious!

6.1.1.Check out the records with very high price

```
df[df.Price>=200]
```

Size \	App	Category	Rating	Reviews
4197	most expensive app (H)	FAMILY	4.3	6
1500.0				
4362	☐ I'm rich	LIFESTYLE	3.8	718
26000.0				
4367	I'm Rich - Trump Edition	LIFESTYLE	3.6	275
7300.0				
5351	I am rich	LIFESTYLE	3.8	3547

1800.0						
5354	I am Rich Plus	FAMILY	4.0	856		
8700.0						
5355	I am rich VIP	LIFESTYLE	3.8	411		
2600.0						
5356	I Am Rich Premium	FINANCE	4.1	1867		
4700.0						
5357	I am extremely Rich	LIFESTYLE	2.9	41		
2900.0						
5358	I am Rich!	FINANCE	3.8	93		
22000.0						
5359	I am rich(premium)	FINANCE	3.5	472		
NaN						
5362	I Am Rich Pro	FAMILY	4.4	201		
2700.0						
5364	I am rich (Most expensive app)	FINANCE	4.1	129		
2700.0						
5366	I Am Rich	FAMILY	3.6	217		
4900.0						
5369	I am Rich	FINANCE	4.3	180		
3800.0						
5373	I AM RICH PRO PLUS	FINANCE	4.0	36		
41000.0						

Updated \	Installs	Type	Price	Content	Rating	Genres	Last
4197	100	Paid	399.99		Everyone	Entertainment	July
16, 2018							
4362	10000	Paid	399.99		Everyone	Lifestyle	March
11, 2018							
4367	10000	Paid	400.00		Everyone	Lifestyle	May
3, 2018							
5351	100000	Paid	399.99		Everyone	Lifestyle	January
12, 2018							
5354	10000	Paid	399.99		Everyone	Entertainment	May
19, 2018							
5355	10000	Paid	299.99		Everyone	Lifestyle	July
21, 2018							
5356	50000	Paid	399.99		Everyone	Finance	November
12, 2017							
5357	1000	Paid	379.99		Everyone	Lifestyle	July
1, 2018							
5358	1000	Paid	399.99		Everyone	Finance	December
11, 2017							
5359	5000	Paid	399.99		Everyone	Finance	May
1, 2017							
5362	5000	Paid	399.99		Everyone	Entertainment	May
30, 2017							
5364	1000	Paid	399.99		Teen	Finance	December
6, 2017							

5366	10000	Paid	389.99	Everyone	Entertainment	June
22, 2018						
5369	5000	Paid	399.99	Everyone	Finance	March
22, 2018						
5373	1000	Paid	399.99	Everyone	Finance	June
25, 2018						

	Current Ver	Android Ver
4197	1.0	7.0 and up
4362	1.0.0	4.4 and up
4367	1.0.1	4.1 and up
5351	2.0	4.0.3 and up
5354	3.0	4.4 and up
5355	1.1.1	4.3 and up
5356	1.6	4.0 and up
5357	1.0	4.0 and up
5358	1.0	4.1 and up
5359	3.4	4.4 and up
5362	1.54	1.6 and up
5364	2	4.0.3 and up
5366	1.5	4.2 and up
5369	1.0	4.2 and up
5373	1.0.2	4.1 and up

```
a=len(df[df.Price>=200])
print(a,"apps are>200 ")
```

15 apps are>200

6.1.2.Drop these as most seem to be junk apps

```
df.drop(df.index[df.Price>=200],inplace=True)
len(df.Price)
```

9078

5.2.Reviews: Very few apps have very high number of reviews. These are all star apps that don't help with the analysis and, in fact, will skew it. Drop records having more than 2 million reviews.

```
df.drop(df.index[df.Reviews>2000000],inplace=True)
len(df.index)
```

8717

6.3.Installs: There seems to be some outliers in this field too. Apps having very high number of installs should be dropped from the analysis.

6.3.1Find out the different percentiles – 10, 25, 50, 70, 90, 95, 99

```
installs_10_perc=np.percentile(df.Installs,10)
print(installs_10_perc)
```

```
1000.0
```

```
installs_25_perc=np.percentile(df.Installs,25)  
installs_25_perc
```

```
10000.0
```

```
installs_50_perc=np.percentile(df.Installs,50)  
installs_50_perc
```

```
500000.0
```

```
installs_75_perc=np.percentile(df.Installs,70)  
installs_50_perc
```

```
500000.0
```

```
installs_90_perc=np.percentile(df.Installs,90)  
installs_90_perc
```

```
10000000.0
```

```
installs_95_perc=np.percentile(df.Installs,95)  
installs_95_perc
```

```
10000000.0
```


```
installs_99_perc=np.percentile(df.Installs,99)  
installs_99_perc
```

```
100000000.0
```

1. Bivariate analysis: Let's look at how the available predictors relate to the variable of interest, i.e., our target variable rating. Make scatter plots (for numeric features) and box plots (for character features) to assess the relations between rating and the other features.

7.1.Make scatter plot/joinplot for Rating vs. Price

```
joint_ch=sns.jointplot(data=df,x='Rating',y='Price')
```




7.1.1.ans.What pattern do you observe? Does rating increase with price?


i observed that the rating and price are not in a good relationship. price has negatively impacted on rating or seems like price has limited impact on rating.

7.2.Make scatter plot/joinplot for Rating vs. Size

```
x=(df['Rating'])
y=(df['Size'])
plt.scatter(x,y)
plt.xlabel('rating')
plt.ylabel('size')
plt.title('raiting vs size scatter plot')
plt.show()
```



```
sns.jointplot(data=df,x='Size',y='Rating')  
<seaborn.axisgrid.JointGrid at 0x7fa239391e50>
```




7.2.1. Are heavier apps rated better?

this plot showing us that most of apps rated between (3.5 to 5.0) almost data distributed evenly. but relationship of this two variable unconvinced.

Make scatter plot/joinplot for Rating vs. Reviews

```
sns.jointplot(data=df,x='Rating',y='Reviews')
```

```
<seaborn.axisgrid.JointGrid at 0x7fa239416cd0>
```




7.3.1. Does more review mean a better rating always?

this plot shown some relation between rating and review. this is left skew plot. high rating having high reviews.

7.4. Make boxplot for Rating vs. Content Rating

```
sns.boxplot(data=df, x='Rating', y='Content Rating')
```

```
<AxesSubplot:xlabel='Rating', ylabel='Content Rating'>
```




7.4.1. Is there any difference **in** the ratings? Are some types liked better?

Adult only 18+ has slightly higher rating. others content rating are seems as same.

7.5. Make boxplot **for** Ratings vs. Category

```
dim_box=(15,12)
plt.subplots(figsize=dim_box)
sns.boxplot(data=df,x='Rating',y='Category')
```

```
<AxesSubplot:xlabel='Rating', ylabel='Category'>
```

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
```

8. Data preprocessing

For the steps below, create a copy of the dataframe to make all the edits. Name it `inp1`.

8.1.Reviews and Install have some values that are still relatively very high. Before building a linear regression model, you need to reduce the skew. Apply log transformation (`np.log1p`) to Reviews and Installs.

```
mod_1=df.copy()
mod_1.Reviews=mod_1.Reviews.apply(np.log1p)
mod_1.Installs=mod_1.Installs.apply(np.log1p)
```

8.2.Drop columns App, Last Updated, Current Ver, and Android Ver. These variables are not useful for our task.

```
mod_1.drop(columns=['App', 'Last Updated', 'Current Ver', 'Android Ver'],inplace=True)
```

```
mod_1.shape
```

```
(7069, 9)
```

8.3.Get dummy columns for Category, Genres, and Content Rating. This needs to be done as the models do not understand categorical data, and all data should be numeric. Dummy

encoding is one way to convert character- fields to numeric. Name of dataframe should be inp2.

```
inp2=pd.get_dummies(mod_1)
```

```
inp2.shape
```

```
(7069, 158)
```

1. Train test split and apply 70-30 split. Name the new dataframes df_train and df_test.

```
data=inp2.drop(columns='Rating')
```

```
data.shape
```

```
(7069, 157)
```

```
target=pd.DataFrame(inp2.Rating)
```

```
target.shape
```

```
(7069, 1)
```

1. Separate the dataframes into X_train, y_train, X_test, and y_test.

```
x_train,x_test,y_train,y_test=train_test_split(data,target,test_size=0.30,random_state=32)
```

```
print("x_train =",x_train.shape)
```

```
print('x_test =',x_test.shape)
```

```
print("y_train =",y_train.shape)
```

```
print('y_test =',y_test.shape)
```

```
x_train = (4948, 157)
```

```
x_test = (2121, 157)
```

```
y_train = (4948, 1)
```

```
y_test = (2121, 1)
```

11. Model building Use linear regression as the technique Report the R2 on the train set

```
regression=LinearRegression()
```

```
regression.fit(x_train,y_train)
```

```
LinearRegression()
```

```
train_pred=regression.predict(x_train)
```

```
print('r2 value of train set is',r2_score(y_train,train_pred))
```

```
r2 value of train set is 0.19371992182920217
```

1. Make predictions on test set and report R2.

```
test_pred=regression.predict(x_test)
```

```
print('r2 value of train set is',r2_score(y_test,test_pred))
```

```
r2 value of train set is 0.14603814020909311
```